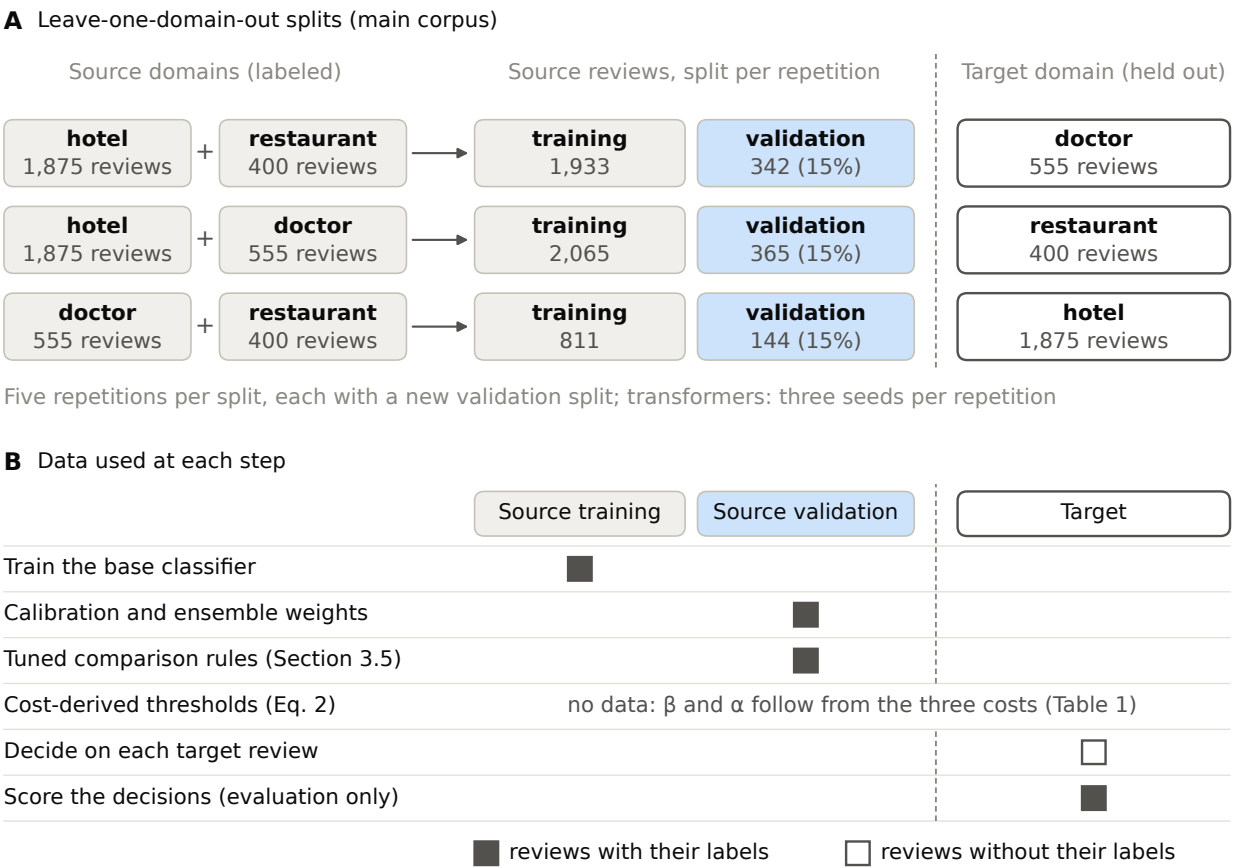
A Leave-one-domain-out splits (main corpus)
Source domains (labeled)
Source reviews, split per repetition
Target domain (held out)
hotel 1,875 reviews + restaurant 400 reviews
training 1,933
validation 342 (15%)
doctor 555 reviews
hotel 1,875 reviews + doctor 555 reviews
training 2,065
validation 365 (15%)
restaurant 400 reviews
doctor 555 reviews + restaurant 400 reviews
training 811
validation 144 (15%)
hotel 1,875 reviews
Five repetitions per split, each with a new validation split; transformers: three seeds per repetition
B Data used at each step
Source training
Source validation
Target
Train the base classifier
Calibration and ensemble weights
Tuned comparison rules (Section 3.5)
Cost-derived thresholds (Eq. 2)
no data: β and α follow from the three costs (Table 1)
Decide on each target review
Score the decisions (evaluation only)
reviews with their labels
reviews without their labels